

SOURCE

May 25, 2024

```
[1]: !pip install pyspark findspark
```

```
Collecting pyspark
  Downloading pyspark-3.5.1.tar.gz (317.0 MB)
    317.0/317.0
MB 2.0 MB/s eta 0:00:00
  Preparing metadata (setup.py) ... done
WARNING: Retrying (Retry(total=4, connect=None, read=None, redirect=None,
status=None)) after connection broken by 'ProtocolError('Connection aborted.',
RemoteDisconnected('Remote end closed connection without response'))':
/simple/findspark/
Collecting findspark
  Downloading findspark-2.0.1-py2.py3-none-any.whl (4.4 kB)
Requirement already satisfied: py4j==0.10.9.7 in /usr/local/lib/python3.10/dist-
packages (from pyspark) (0.10.9.7)
Building wheels for collected packages: pyspark
  Building wheel for pyspark (setup.py) ... done
  Created wheel for pyspark: filename=pyspark-3.5.1-py2.py3-none-any.whl
size=317488491
sha256=b7ca25f273628e79e35111358be0d9d82b370c8b7b894f2fab962a4a8270c6a5
  Stored in directory: /root/.cache/pip/wheels/80/1d/60/2c256ed38dddce2fdd93be54
5214a63e02fbd8d74fb0b7f3a6
Successfully built pyspark
Installing collected packages: findspark, pyspark
Successfully installed findspark-2.0.1 pyspark-3.5.1
```

```
[2]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[3]: mnist_mini = "/content/drive/MyDrive/Colab Notebooks/MMDS/final/mnist_mini.csv"
```

```
[4]: ratings2k = "/content/drive/MyDrive/Colab Notebooks/MMDS/final/ratings2k.csv"
```

```
[5]: stockHVN2022 = "/content/drive/MyDrive/Colab Notebooks/MMDS/final/stockHVN2022.  
      ↪CSV"
```

```
[6]: from pyspark.sql import SparkSession  
      from pyspark.ml.feature import VectorAssembler  
      from pyspark.ml.clustering import KMeans  
      from pyspark.ml.linalg import Vectors as VectorsML  
  
      from pyspark.sql import Row, DataFrame  
      from pyspark.sql.functions import udf, col, sum, count, ↪  
          ↪monotonically_increasing_id, when, to_date, month, ↪  
              lag, col, month, array  
  
      from pyspark.sql.types import IntegerType, FloatType, ArrayType, DoubleType  
      from pyspark.mllib.linalg.distributed import RowMatrix  
      from pyspark.mllib.linalg import DenseMatrix, Vectors  
      from pyspark.ml.regression import LinearRegression  
      from pyspark.ml.classification import MultilayerPerceptronClassifier, ↪  
          ↪RandomForestClassifier, LinearSVC, OneVsRest  
      from pyspark.ml.evaluation import MulticlassClassificationEvaluator, ↪  
          ↪RegressionEvaluator  
      from pyspark.sql.window import Window  
      from pyspark.ml.recommendation import ALS  
  
      import math  
      import matplotlib.pyplot as plt  
      import numpy as np  
      from mpl_toolkits.mplot3d import Axes3D
```

```
[7]: spark = SparkSession.builder.appName("Final_Project").getOrCreate()
```

1 Task 1: Clustering

```
[8]: class Clustering:  
      def __init__(self, data_path:str ,k:int, weighted_rows:list):  
          self.data = None  
          self.sum_count_df = None  
          self.data_path = data_path  
          self.k = k  
          self.weighted_rows = weighted_rows  
          self.transformed_data = None  
  
      def load_data(self):  
          self.data = spark.read.csv(self.data_path, header=False, ↪  
          ↪inferSchema=True)
```

```

def preprocess_data(self):
    assembler = VectorAssembler(inputCols=self.data.columns[1:],
    ↪outputCol="features")
    self.data = assembler.transform(self.data).select("_c0", "features")

    self.data = self.data.withColumn("row_index",
    ↪monotonically_increasing_id())
    self.data = self.data.withColumn("weight", when(self.data["row_index"]
    ↪isin(self.weighted_rows), 100).otherwise(1)) \
        .drop("row_index")

def clustering(self):
    kmeans = KMeans(featuresCol="features", k=self.k, weightCol="weight",
    ↪seed=42)
    model = kmeans.fit(self.data)
    transformed_data = model.transform(self.data)
    self.transformed_data = transformed_data

    l_clusters = model.clusterCenters()
    df_centers = spark.createDataFrame([(int(i), VectorsML.
    ↪dense(l_clusters[i]))
        for i in range(len(l_clusters))], ["prediction", "center"])

    transformed_data = transformed_data.join(df_centers, on='prediction',
    ↪how='left')

    get_dist = udf(lambda features, center: float(math.sqrt(VectorsML.
    ↪squared_distance(features, center))), FloatType())
    transformed_data = transformed_data.withColumn('dist',
    ↪get_dist(col('features'), col('center')))

    self.sum_count_df = transformed_data.groupBy("prediction") \
        .agg(sum("dist").alias("sum_distance"),
            count("*").alias("count")) \
        .withColumn('average_dist', col('sum_distance') / col('count')) \

def plot_average_distance(self):
    result_list = self.sum_count_df.select("prediction", "average_dist").
    ↪collect()

    prediction = [row.prediction for row in result_list]
    average_dist = [row.average_dist for row in result_list]

    fig, ax = plt.subplots()
    ax.bar(prediction, average_dist)
    ax.set_xlabel('k')

```

```

ax.set_ylabel('Average Distance')
ax.set_title('Average Distance by Prediction')
plt.xticks(prediction)
plt.show()

```

```

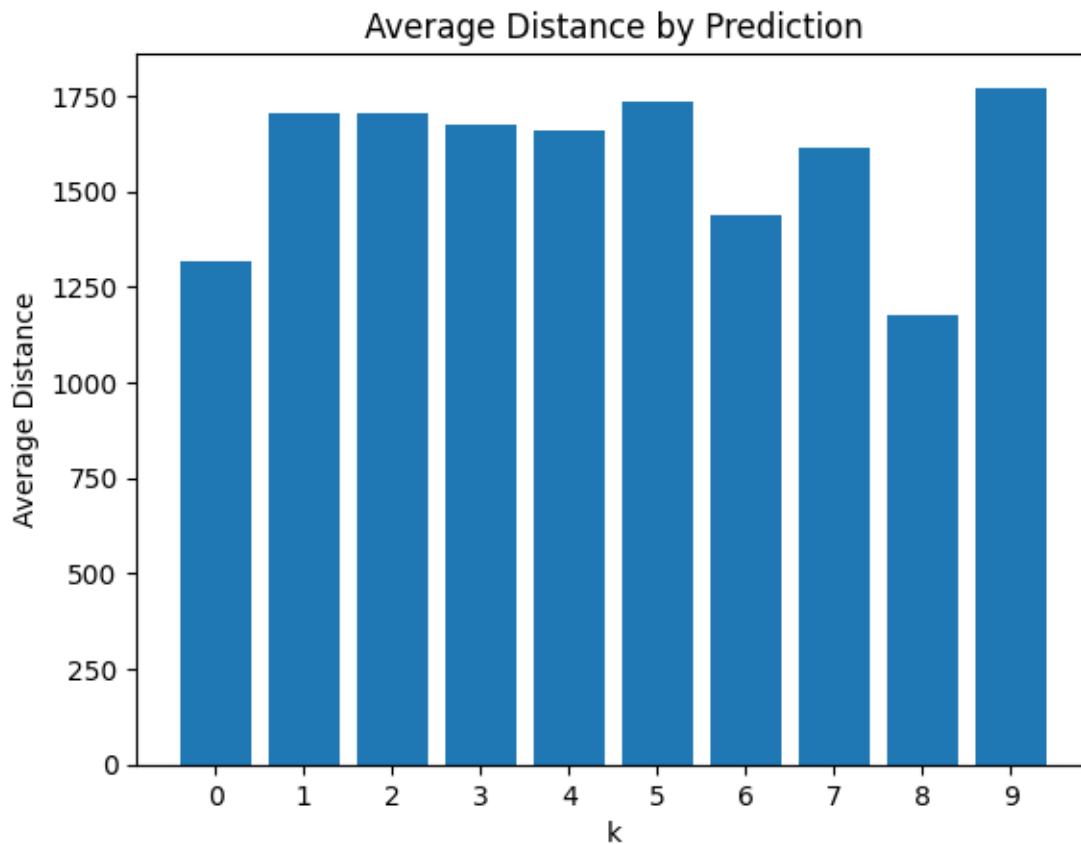
def run(self):
    self.load_data()
    self.preprocess_data()
    self.clustering()
    self.plot_average_distance()

```

```

[9]: if __name__ == "__main__":
    weighted_rows = [0, 1, 2, 3, 4, 7, 8, 11, 18, 61]
    clustering = Clustering(data_path = mnist_mini, k = 10, weighted_rows = weighted_rows)
    clustering.run()

```



2 Task 2: Dimensionality Deduction with SVD

```
[10]: class SVD:
    def __init__(self, data_path:str, clustering_result:DataFrame):
        self.data = None
        self.data_path = data_path
        self.processed_data = None
        self.clustering_result = clustering_result

    def load_data(self):
        self.data = spark.read.csv(self.data_path, header=False,
        ↪inferSchema=True)

    def preprocess_data(self):
        assembler = VectorAssembler(inputCols=self.data.columns[1:],
        ↪outputCol="features")
        self.data = assembler.transform(self.data) \
            .select("_c0", "features")

    def svd(self):
        ml_vectors = self.data.select("features").rdd.map(lambda row: Vectors.
        ↪dense(row.features.toArray()))
        row=RowMatrix(ml_vectors)
        svd = row.computeSVD(3, True)
        u = svd.U
        s = DenseMatrix(len(svd.s), len(svd.s), np.diag(svd.s).ravel("F"))

    def row_to_dict(row):
        return Row(features=Vectors.dense(row.toArray()))

    row_matrix_rows = u.multiply(s).rows.map(row_to_dict)
    df = spark.createDataFrame(row_matrix_rows)

    vector_to_array = udf(lambda v: v.toArray().tolist(),
    ↪ArrayType(DoubleType()))
    df = df.withColumn("features_array", vector_to_array(df["features"]))
    df = df.select([df["features_array"][i].alias(f"_c{i+1}") for i in
    ↪range(3)])

    data = self.clustering_result.withColumn("unique_id",
    ↪monotonically_increasing_id())
    df = df.withColumn("unique_id", monotonically_increasing_id())
    merged_df = data[['_c0', 'unique_id', "prediction"]].join(df,
    ↪"unique_id").drop("unique_id")

    self.processed_data = merged_df.sample(fraction=0.01).limit(100)
```

```

def plot_3d(self):
    color = ["red", "blue", "orange", "green", "purple", "brown", "pink", "yellow", "black", "cyan"]
    marker = ["v", "*", ".", ",", "X", "P", "s", "D", "^", "v"]

    fig = plt.figure()
    ax = fig.add_subplot(111, projection='3d')

    for i in range(0, 10):
        data = self.processed_data.filter(self.processed_data.prediction == i)

        x = data.select('_c1').rdd.map(lambda row: row['_c1']).collect()
        y = data.select('_c2').rdd.map(lambda row: row['_c2']).collect()
        z = data.select('_c3').rdd.map(lambda row: row['_c3']).collect()
        ax.scatter(x, y, z, c=color[i], marker=marker[i])

    ax.set_title('3D K-Means Clustering with k=10 (100 Random Datapoint)')
    ax.set_xlabel('Feature 1')
    ax.set_ylabel('Feature 2')
    ax.set_zlabel('Feature 3')
    plt.show()

def run(self):
    self.load_data()
    self.preprocess_data()
    self.svd()
    self.plot_3d()


```

```

[11]: if __name__ == "__main__":
        svd = SVD(data_path = mnist_mini, clustering_result = clustering.
        transformed_data)
        svd.run()

```

3D K-Means Clustering with k=10 (100 Random Datapoint)



3 Task 3: Recommendation with Collaborative Filtering

```
[12]: class CollaborativeFiltering:
    def __init__(self, data_path: str = "ratings2k.csv", num_user_blocks_range:
        array = np.arange(10, 21)):
        self.data = None
        self.train_data = None
        self.test_data = None
        self.als_model = None
        self.data_path = data_path
        self.num_user_blocks_range = num_user_blocks_range

    def load_data(self):
        self.data = spark.read.csv(self.data_path, header=True,
        inferSchema=True)

    def split_data(self):
        self.train_data, self.test_data = self.data.randomSplit([0.7, 0.3],
        seed=42)

    def train_model(self, rank=10, max_iter=20, reg_param=0.01):
```

```

        als = ALS(rank=rank, maxIter=max_iter, regParam=reg_param,
↪userCol="user", itemCol="item", ratingCol="rating",
            coldStartStrategy="drop", nonnegative=True)

        self.als_model = als.fit(self.train_data)

    def evaluate_model(self):
        predictions = self.als_model.transform(self.test_data)
        evaluator = RegressionEvaluator(metricName="mse", labelCol="rating",
↪predictionCol="prediction")
        mse = evaluator.evaluate(predictions)
        return mse

    def visualize_results(self):
        mse_values = []
        ranks = np.arange(10, 21)
        for rank in ranks:
            self.train_model(rank = rank)
            mse = self.evaluate_model()
            mse_values.append(mse)

        plt.figure(figsize=(10, 6))
        plt.bar(np.arange(10, 21), mse_values)
        plt.xlabel("The number of similar User")
        plt.ylabel("MeanSquare Error (MSE)")
        plt.title("MSE vs Number of similar users")
        plt.xticks(np.arange(10, 21))
        plt.show()

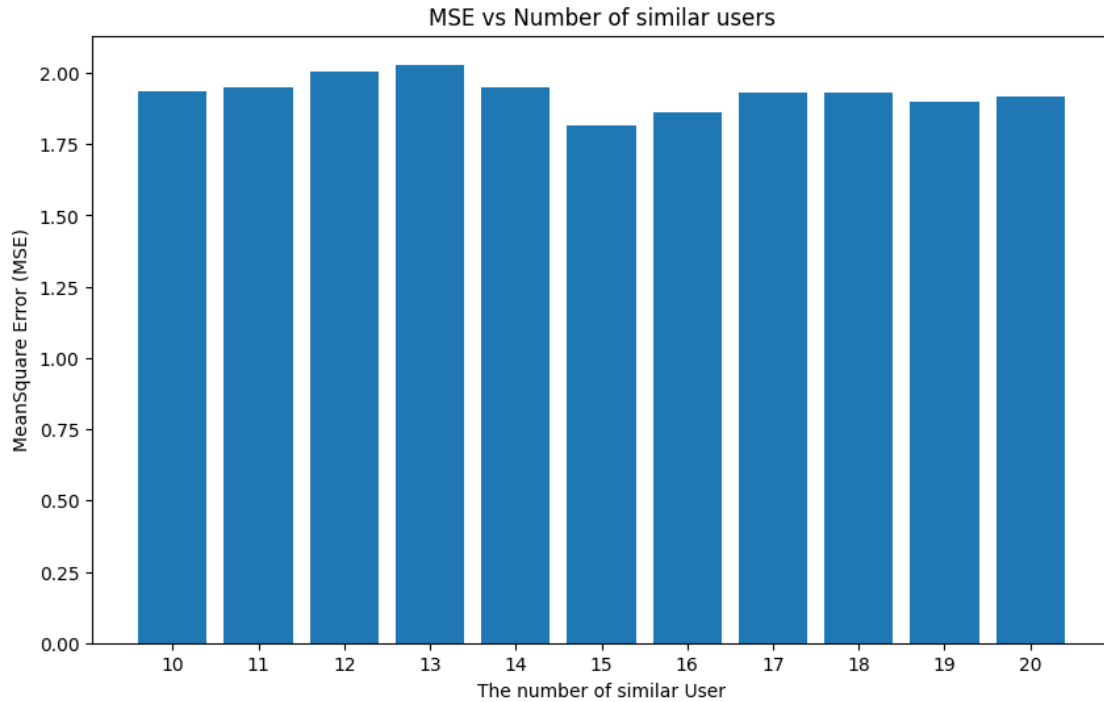
    def run(self):
        self.load_data()
        self.split_data()
        self.visualize_results()

```

```

[13]: if __name__ == "__main__":
        collaborativeFiltering = CollaborativeFiltering(data_path = ratings2k)
        collaborativeFiltering.run()

```

4 Task 4: Stock price regression

```
[14]: class StockPriceRegression:
    def __init__(self, k_prev_dates:int = 5, data_path:str = "stockHVN2022.
    ↪csv"):
        self.k_prev_dates = k_prev_dates
        self.data = None
        self.data_path = data_path
        self.train_data, self.test_data = None, None
        self.model = None
        self.train_mse, self.test_mse = None, None

    def load_data(self):
        self.data = spark.read.csv(self.data_path, header=True,
    ↪inferSchema=True)

    def preprocess_data(self):
        self.data = self.data.withColumn("Ngay", to_date(self.data.Ngay, "dd/MM/
    ↪yyyy")) \
            .orderBy('Ngay')\
            .withColumn('fluctuation', (col('HVN') - lag('HVN', 1).
    ↪over(Window.orderBy('Ngay')) / lag('HVN', 1).over(Window.orderBy('Ngay')))\
            .na.fill(0.0, subset=['fluctuation']))
```

```

        for i in range(1, self.k_prev_dates + 1):
            self.data = self.data.withColumn(f'prev_range_{i}',
            ↪lag('fluctuation', i).over(Window.partitionBy().orderBy('Ngay'))))
            self.data = self.data.dropna()

        train_data = self.data.filter(month('Ngay') <= 6) \
            .drop("Ngay", "HVN")
        test_data = self.data.filter(month('Ngay') > 6) \
            .drop("Ngay", "HVN")

        assembler = VectorAssembler(inputCols=list(train_data.columns[:-1])[:
        ↪-1]), outputCol='features')
        self.train_data = assembler.transform(train_data)
        self.test_data = assembler.transform(test_data)

    def train_model(self):
        lr = LinearRegression(featuresCol='features', labelCol='fluctuation')
        self.model = lr.fit(self.train_data)

    def evaluate_model(self):
        evaluator = RegressionEvaluator(predictionCol='prediction',
        ↪labelCol='fluctuation', metricName='mse')


        self.train_mse = evaluator.evaluate(self.model.transform(self.
        ↪train_data))
        self.test_mse = evaluator.evaluate(self.model.transform(self.test_data))

    def plot_results(self):
        plt.figure(figsize=(8, 6))
        plt.bar(['Training', 'Test'], [self.train_mse, self.test_mse],
        ↪color=['blue', 'green'])
        plt.xlabel('Dataset')
        plt.ylabel('Mean Squared Error')
        plt.title('Linear Regression Model Performance')
        for i, v in enumerate([self.train_mse, self.test_mse]):
            plt.text(i, v, str(v), ha='center', va='bottom')
        plt.show()

    def run(self):
        self.load_data()
        self.preprocess_data()
        self.train_model()
        self.evaluate_model()
        self.plot_results()

```

```
[15]: if __name__ == "__main__":
    stockPriceRegression = StockPriceRegression(data_path = stockHVN2022)
    stockPriceRegression.run()
```



5 Task 5: Multi-class classification

```
[16]: class MultiClassClassification:
    def __init__(self, data_path:str):
        self.data = None
        self.classification = ["MLP", "RandomForest", "LinearSVC"]
        self.train_accuracies = []
        self.test_accuracies = []
        self.data_path = data_path

    def load_data(self):
        self.data = spark.read.csv(self.data_path, header=False,
        inferSchema=True)

    def preprocess_data(self):
```

```

        assembler = VectorAssembler(inputCols=self.data.columns[1:],
        ↪outputCol="features")
        self.data = assembler.transform(self.data) \
            .select("_c0", "features") \
            .withColumnRenamed("_c0", "label")

    def run_classification(self):
        (train_data, test_data) = self.data.randomSplit([0.7, 0.3], seed=42)

        for classifier_name in self.classification:
            if classifier_name == "MLP":
                classifier = MultilayerPerceptronClassifier(maxIter=100,
        ↪layers=[784, 128, 64, 10], blockSize=128, seed=42)
            elif classifier_name == "RandomForest":
                classifier = RandomForestClassifier(numTrees=100, maxDepth=10,
        ↪seed=42)
            else:
                classifier = OneVsRest(classifier=LinearSVC(maxIter=100,
        ↪regParam=0.1))

            model = classifier.fit(train_data)

            evaluator = MulticlassClassificationEvaluator(metricName="accuracy")
            train_accuracy = evaluator.evaluate(model.transform(train_data))
            test_accuracy = evaluator.evaluate(model.transform(test_data))

            self.train accuracies.append(train_accuracy)
            self.test accuracies.append(test_accuracy)

    def plot accuracies(self):
        x = np.arange(len(self.classification))
        fig, ax = plt.subplots()
        bar_width = 0.35

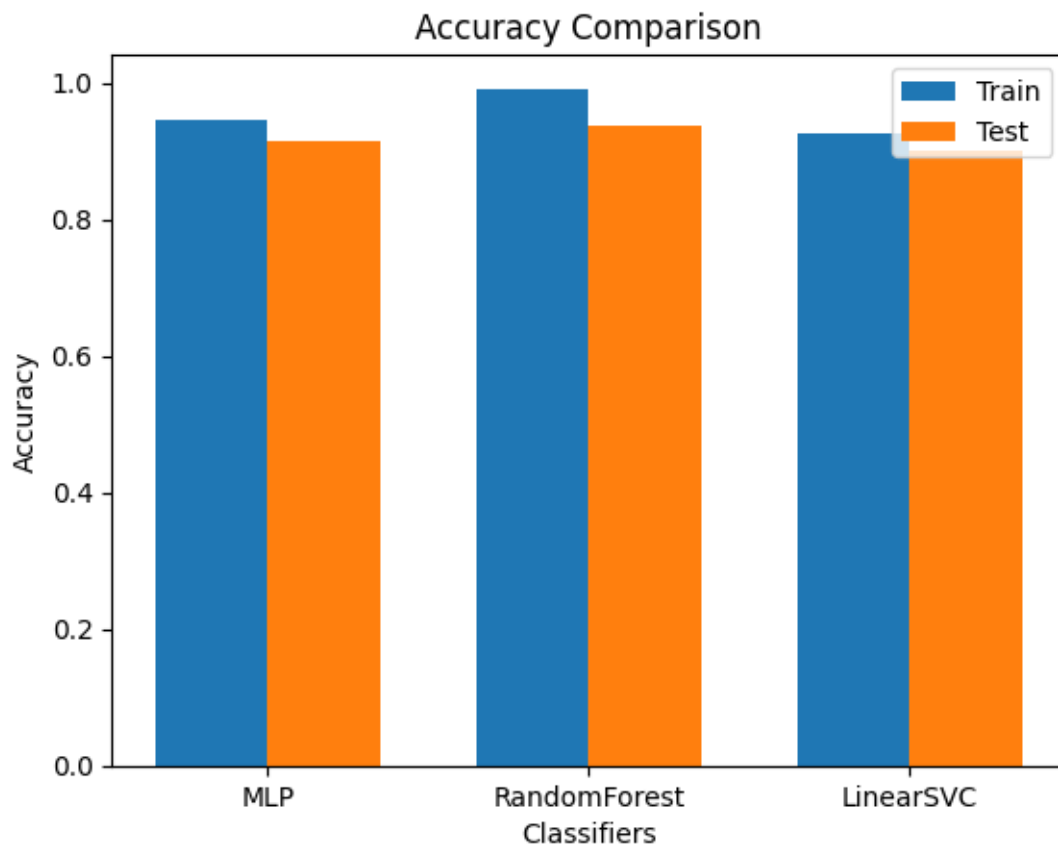
        rects1 = ax.bar(x - bar_width / 2, self.train accuracies, bar_width,
        ↪label='Train')
        rects2 = ax.bar(x + bar_width / 2, self.test accuracies, bar_width,
        ↪label='Test')

        ax.set_xlabel('Classifiers')
        ax.set_ylabel('Accuracy')
        ax.set_title('Accuracy Comparison')
        ax.set_xticks(x)
        ax.set_xticklabels(self.classification)
        ax.legend()
        plt.show()

```

```
def run(self):  
    self.load_data()  
    self.preprocess_data()  
    self.run_classification()  
    self.plot accuracies()
```

```
[17]: if __name__ == "__main__":  
    multiClassClassification = MultiClassClassification(data_path = mnist_mini)  
    multiClassClassification.run()
```



```
[18]: spark.stop()
```